

Paper Reviews

Introduction

This repository is a collection of short reviews of papers in Deep Learning. Reviews are based on papers which explore novel groundbreaking ideas or consist of theoretically rich concepts. A total of 2 short reviews are typically added each week which explore essential aspects of the work, its technical innovation and new questions and ideas raised by the work. Length of each review is 1 page excluding references. Each review is based on a fixed set of guidelines which are given [here](#). If you would like to contribute to the reviews then please read [this](#).

Available Paper Reviews

Review Number	Paper Title	Author List	Review Link	Contributor
1	Action and Perception as Divergence Minimization	Danijar Hafner, Pedro A. Ortega, Jimmy Ba, Thomas Parr, Karl Friston, Nicolas Heess	TBA	@karush17
2	Momentum Contrast for Unsupervised Visual Representation Learning	Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, Ross Girshick	TBA	@karush17

Review Guildines

This section outlines the guidelines which are used to write reviews for the repository. Note that these guidelines must be strictly followed for providing high quality reviews to the reader.

Paper Selection

Since Deep Learning is a fast-moving field, papers can span a broad variety of topics. In order to shorten the range of literature, papers must be selected using the following rules-

- Any review paper, survey or a long articles should not be considered since these are themselves a review of previous works and reviewing them would defeat the purpose of literature writing.
- Short paper, Journal papers and theoretical works are suitable as these present a single idea which may be of interest to the reader.
- Papers containing experimental results balanced with theory are encouraged as these validate the practical applications of the proposed methods.
- There is no constraint on the publication date and time of papers. However, papers which date back to the 'AI Winter' are highly encouraged since these provide insights which were never looked at for a very long time.
- Papers containing simply applications of pre-existing methods to various regimes are not preferred since they do not provide any new insights into the algorithm itself and only deal with its applicability.
- Workshop papers and incomplete works are also welcome as researchers can always build on these ideas.

- While there is no specific restriction on the content of papers, reviewer must consider fields in AI which are growing and require a new outlook, eg. Reinforcement Learning, Gradient-Free methods, Meta-Learning, Natural Language at Scale, Explainability, etc.

Introduction

This section provides points on writing a good introduction for the review. Note that these review points are not strict and only serve as a guiding principle for drafting a good introduction.

- The introduction should be a high-level idea of the paper and what the work deals with.
- The reviewer must refrain from going into any technical detail and highlight the broad idea of the work and its scope.
- No mathematical terms, definitions, technical explanations or algorithmic details should be provided to the reader.
- The main focus should be on the problem statement and how the method aims to solve this.
- A possible checklist of questions one might want to answer while writing a good introduction is as followed-
 - ☐ What is the main area of the work?
 - ☐ What are some of the open problems in this area?
 - ☐ What has been done so far to address these problems?
 - ☐ What still remains as the scope of this paper?
 - ☐ What does the paper propose?
 - ☐ How does the paper solves this problem?
 - ☐ What theoretical/practical insights does the work achieve?
 - ☐ What could have been (or is) done better to solve the problem?
- Key takeaway- Detail is your enemy!

Methodology

This section deals with the proposed method and its essential aspects in solving the problem highlighted in the introduction section. Following points serve as a guide to writing this section-

- The content must provide a brief overview of the method. This informs the reader about what he is getting into.
- Once an overview has been provided, the draft can start diving into the detail which should be highlighted intuitively.
- Mathematical details must be followed by words, complicated terminology must be explained intuitively using equations or instances from work.
- Reasons related to technical details and their usage must be provided to the reader. The whole point of reading a review is to crisply go over the details of the paper without wanting to read the entire text.
- While the draft should highlight the method and its details, it should also provide the reader with insights from the reviewer's point of view. These could consist of specific reasons for selecting a set of parameter values, usage of a specific technique existing in literature, novel contribution and the reason behind its usage and any improvements/changes from previous works.
- A possible checklist of constituents of this section is as followed-

Critical Analysis

-

New Ideas/Questions

-

Conclusion

-

Contributions

"No one can whistle a symphony. It takes a whole orchestra to play it." – H.E. Luccock

Collaboration is most welcome for this repository. If you believe that you would like to practice your research writing skills or would like to do a literature survey for your independent study, then feel free to hop in!

Reviews in this repository follow a general LaTeX templated which is available in the template folder.

Furthermore, reviews written for the repository must follow the review guidelines which are available [here](#). Feel free to submit your review by opening a pull request.